



EXPLORING IMAGE CLASSIFICATION WITH THE CIFAR-10 DATASET



TEAM 4

1. SHUBHAM SINGH
2. A. AKHILA
3. CH. RAHUL
4. A. HARI HARAN

DESCRIPTION:

The CIFAR-10 data consists of 60,000 32x32 color images in 10 classes, with 6000 images per class. There are 50,000 training images and 10,000 test images in the official data. We have preserved the train/test split from the original dataset.

OBJECTIVE:

The objective is to determine the class to which the image belong to using CNN(Convulational Neural Network)

CLASSES:

airplane	cat	deer	
automobile	dog	bird	truck
ship	horse	frog	

Data Preparation:

DATA VISUALIZATION

truck



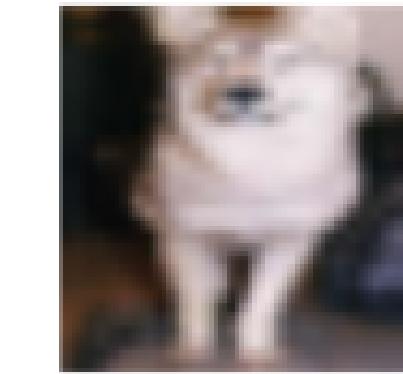
ship



airplane



dog



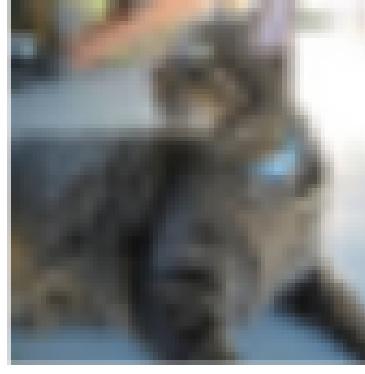
frog



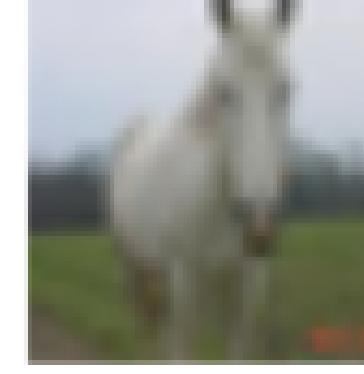
truck



cat



horse



truck



deer



truck



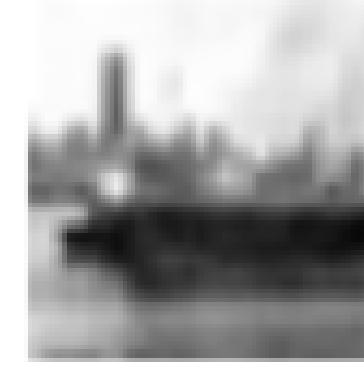
frog



horse



ship



frog



horse



ship



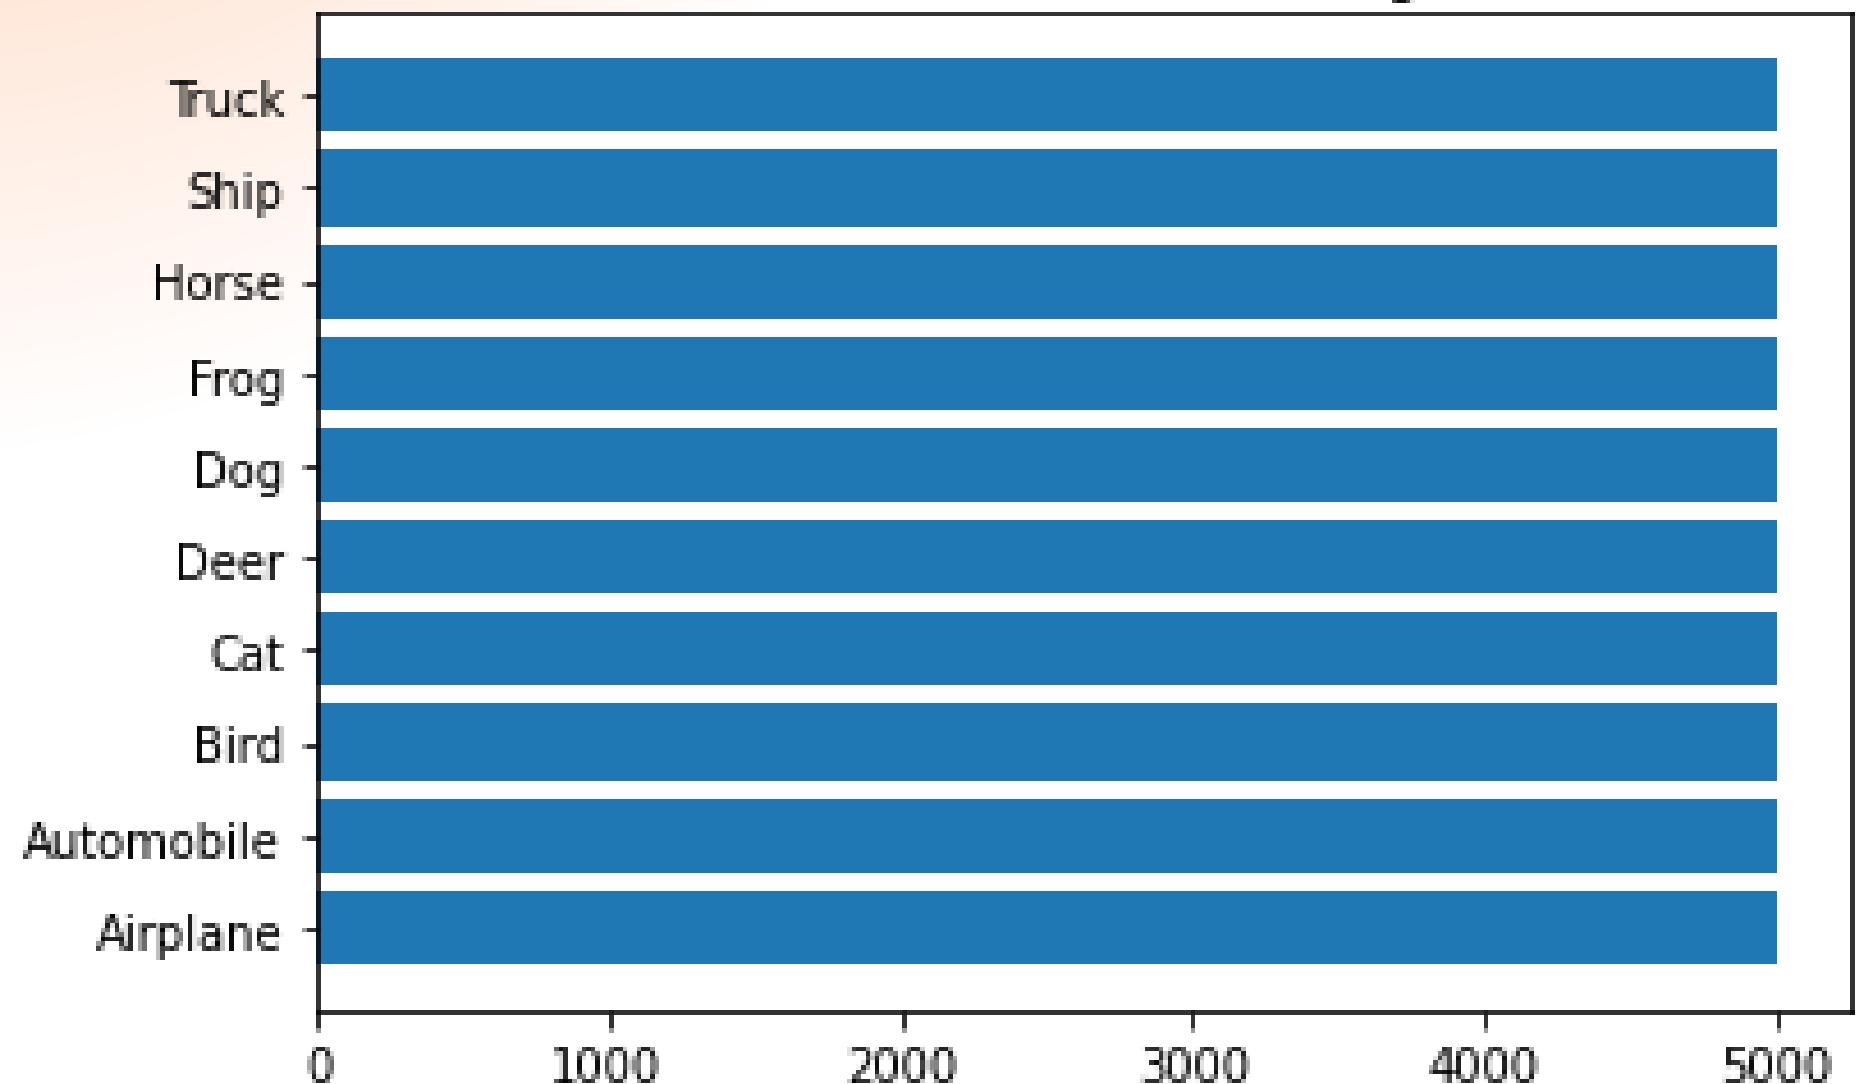
cat



Shape of the image - 32x32 pixels

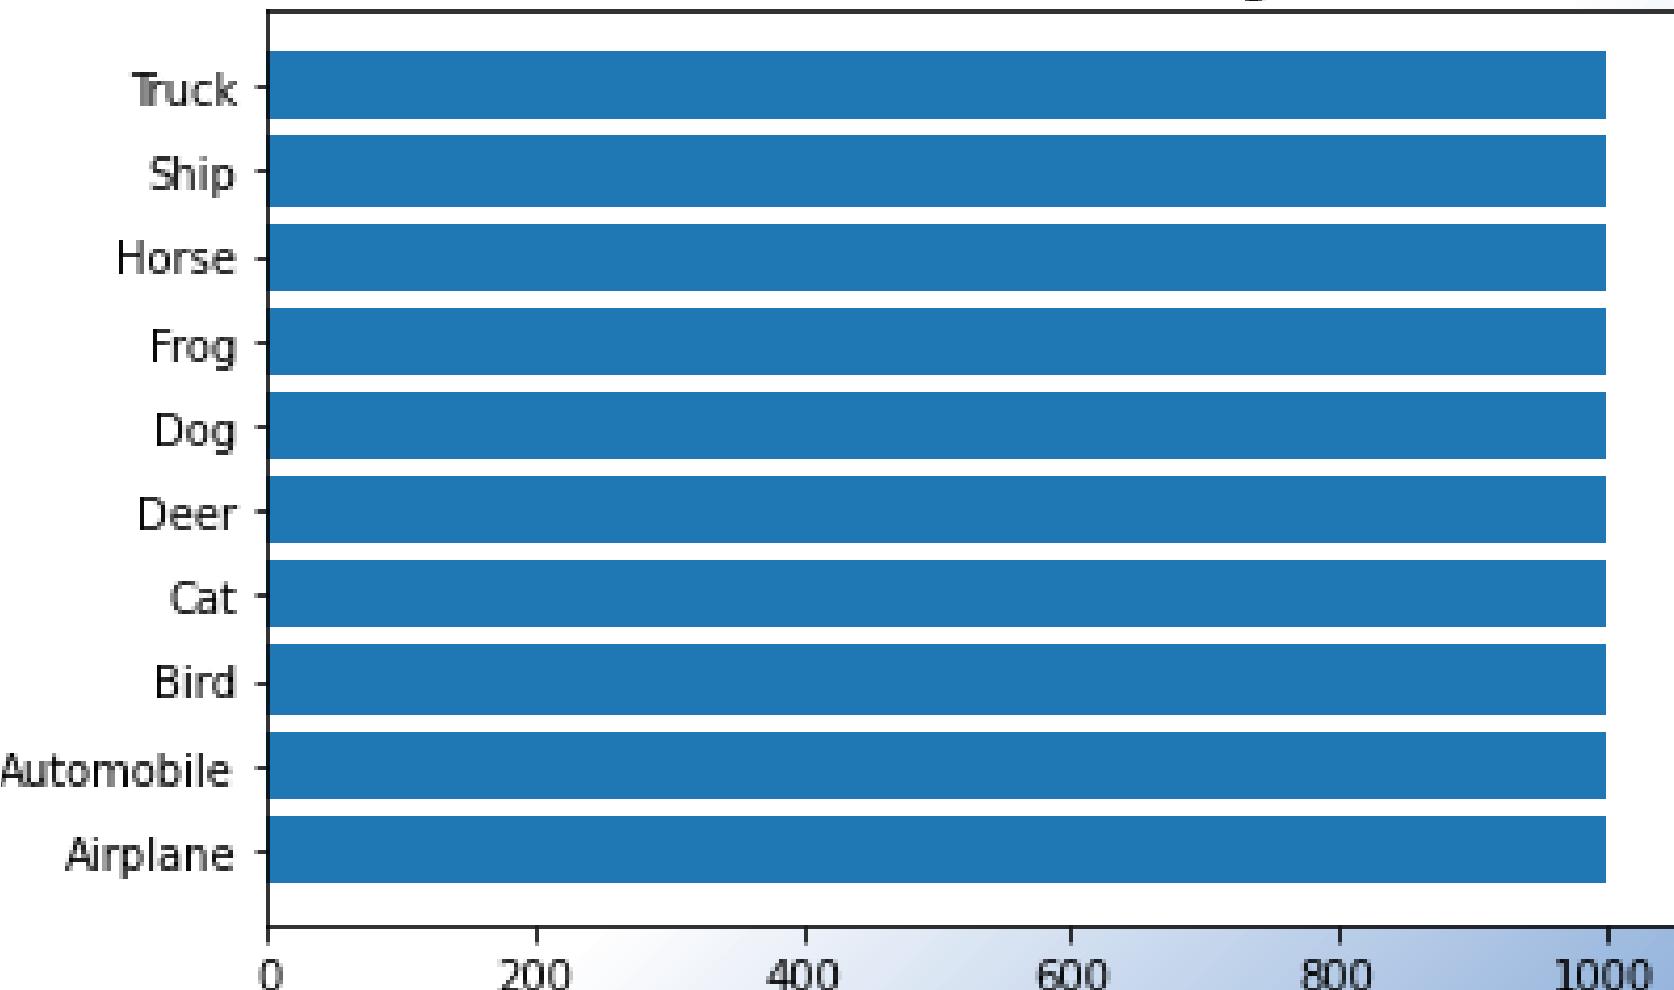
TRAINING DATA

Class distribution in training set



TESTING DATA

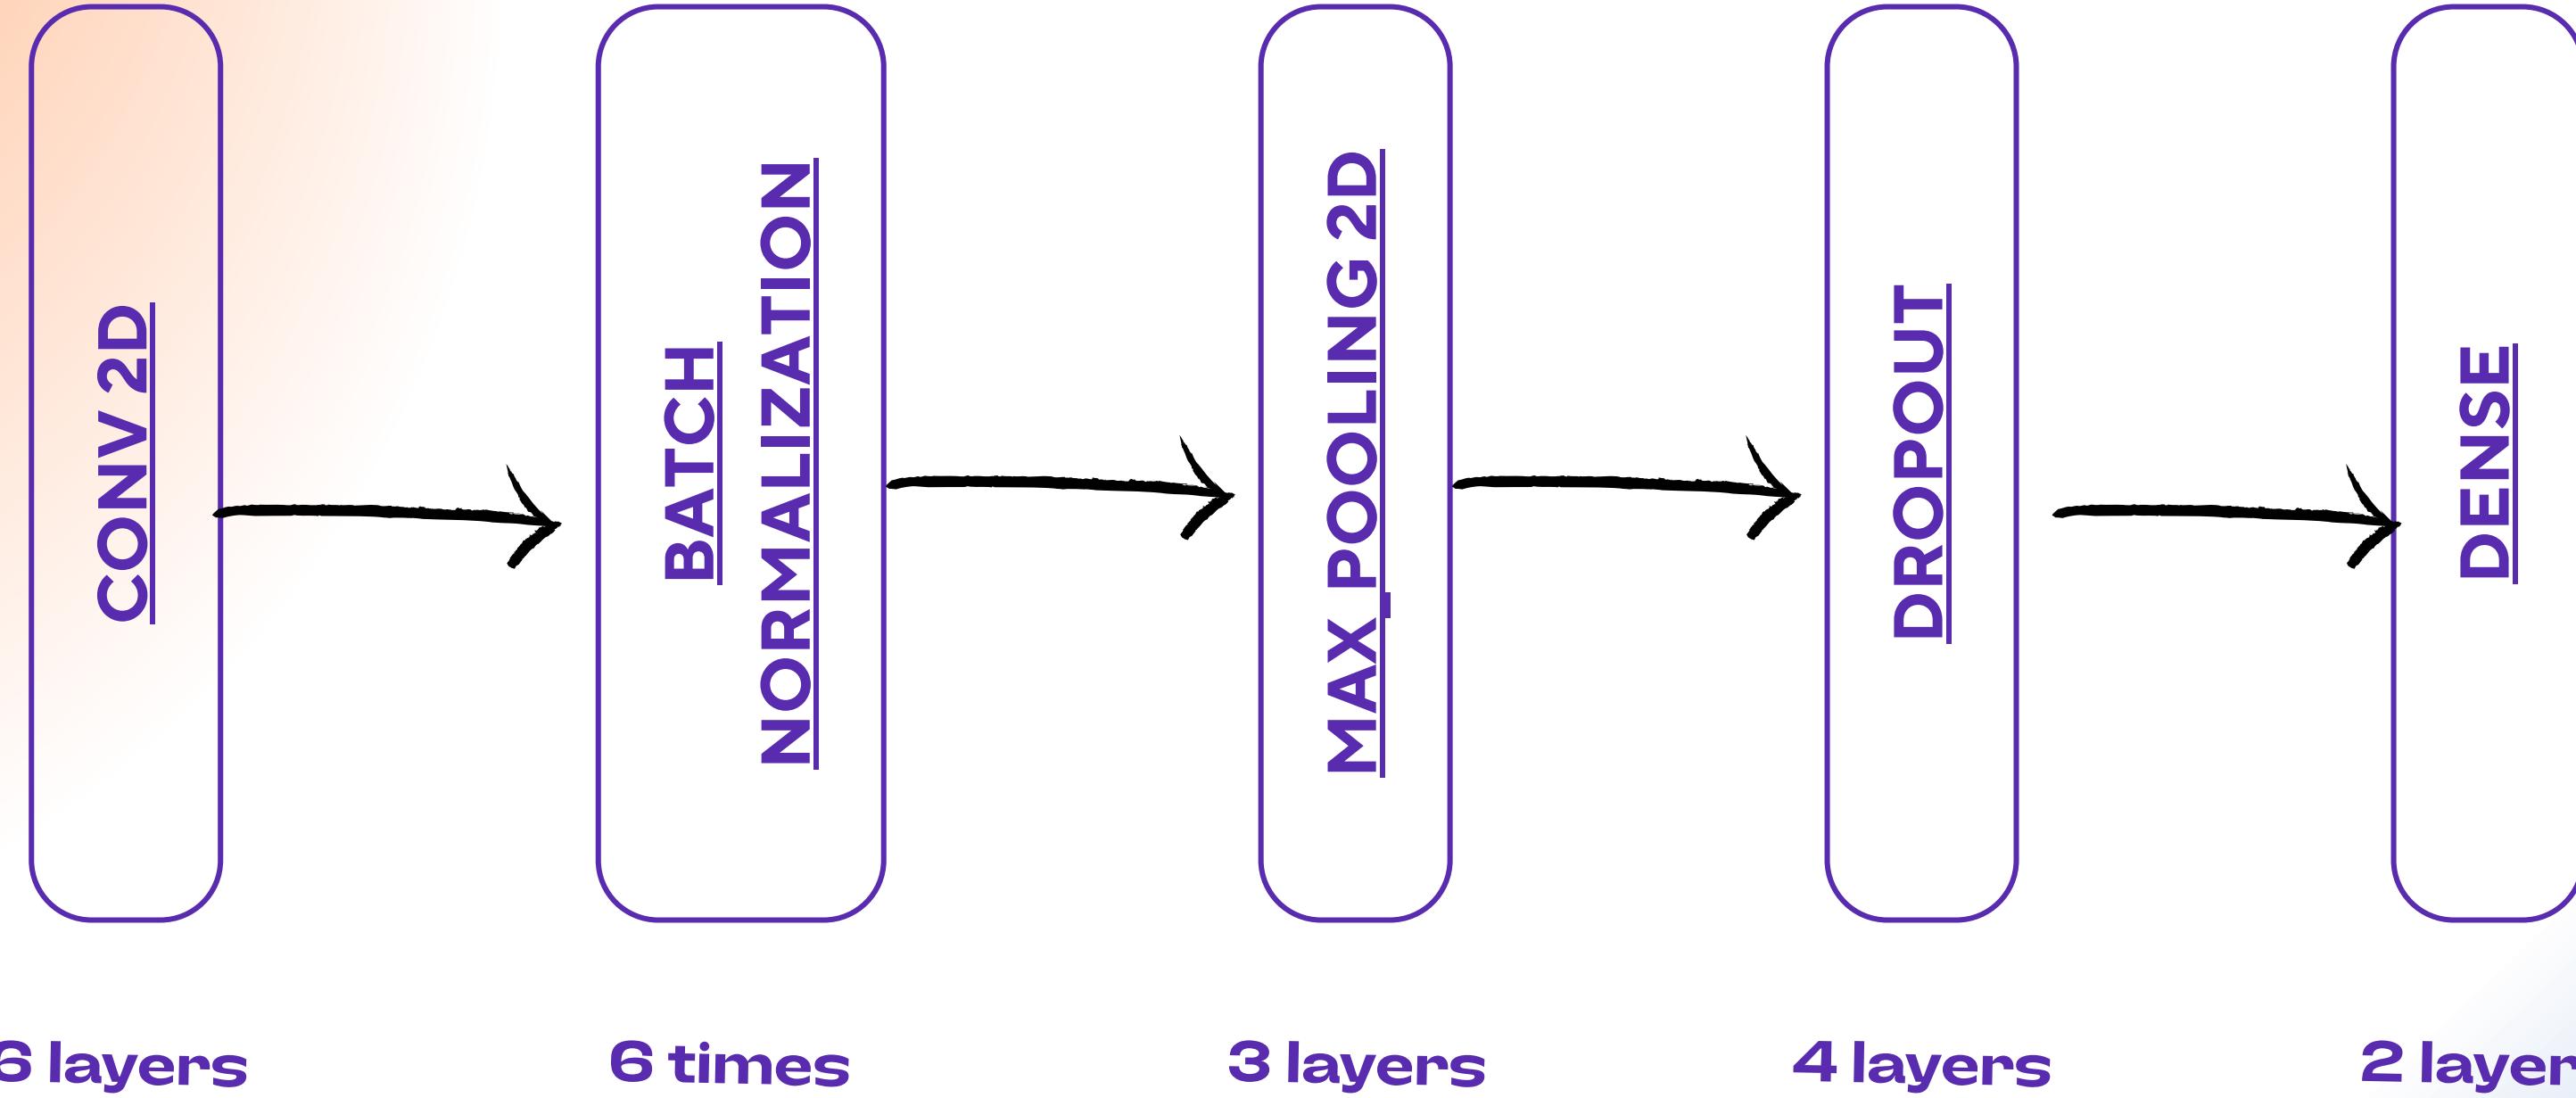
Class distribution in testing set



Parameters:

Optimizers	Adam	SGD
Loss	binary cross entropy	categorical cross entropy
Epochs	min- 20	max - 300
Batch size	min-32	max-64
Metrics	<i>Accuracy</i> <i>Precision</i>	
	<i>Recall</i>	

CNN Summary:



Total params	Trainable params	non trainable params	activation
552362	551466	896	relu,softmax

batch_size	Epoch	optimizer	Loss_function	Accuracy	Precision	Recall	Loss_val
-------------------	--------------	------------------	----------------------	-----------------	------------------	---------------	-----------------

Adam

64	10	Adam	categorical_crossentropy	0.8999	0.9244	0.879	0.394
46	60	Adam	binary_crossentropy	0.9008	0.9205	0.8864	0.0523
64	50	Adam	binary_crossentropy	0.9176	0.9344	0.9026	0.3358
64	300	adam	binary_crossentropy	0.9418	0.9515	0.9351	0.0328

SGD

40	90	SGD	binary_crossentropy	0.8057	0.8466	0.7744	0.0869
50	60	SGD	categorical_crossentropy	0.8396	0.8864	0.7996	0.4622
32	60	SGD	categorical_crossentropy	0.8508	0.8913	0.8144	0.4323

batch_size	Epoch	optimizer	Loss_function	Accuracy	Precision	Recall	Loss_val
-------------------	--------------	------------------	----------------------	-----------------	------------------	---------------	-----------------

binary cross entropy

64	60	Adam	binary_crossentropy	0.8924	0.9142	0.8763	0.0565
46	60	Adam	binary_crossentropy	0.9008	0.9205	0.8864	0.0523
64	50	Adam	binary_crossentropy	0.9176	0.9344	0.9026	0.3358
64	300	adam	binary_crossentropy	0.9418	0.9515	0.9351	0.0328

Categorical cross entropy

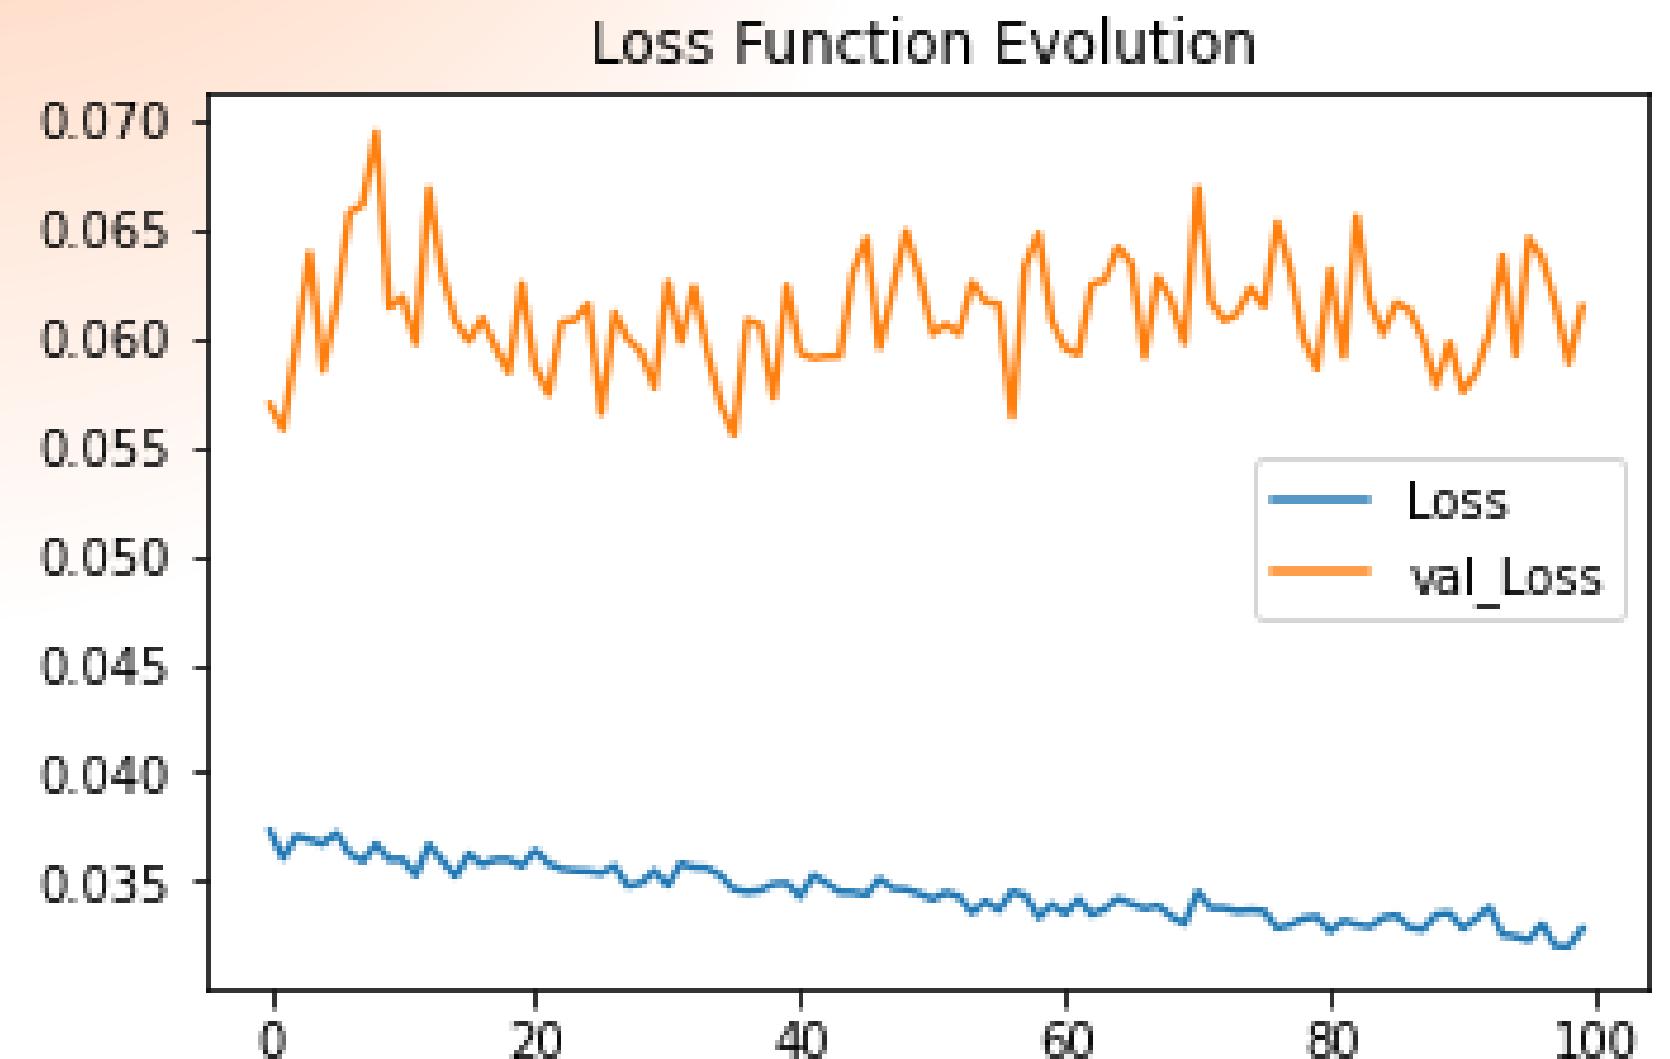
40	40	adam	categorical_crossentropy	0.8963	0.9218	0.8739	0.3744
50	60	Adam	categorical_crossentropy	0.8986	0.9232	0.878	0.2932
64	10	Adam	categorical_crossentropy	0.8999	0.9244	0.879	0.394

batch_size	Epoch	optimizer	Loss_function	Accuracy	Precision	Recall	Loss_val
64	10	Adam	categorical_crossentropy	0.8999	0.9244	0.879	0.394
46	60	Adam	binary_crossentropy	0.9008	0.9205	0.8864	0.0523
64	50	Adam	binary_crossentropy	0.9176	0.9344	0.9026	0.3358
64	300	Adam	binary_crossentropy	0.9418	0.9515	0.9351	0.0328

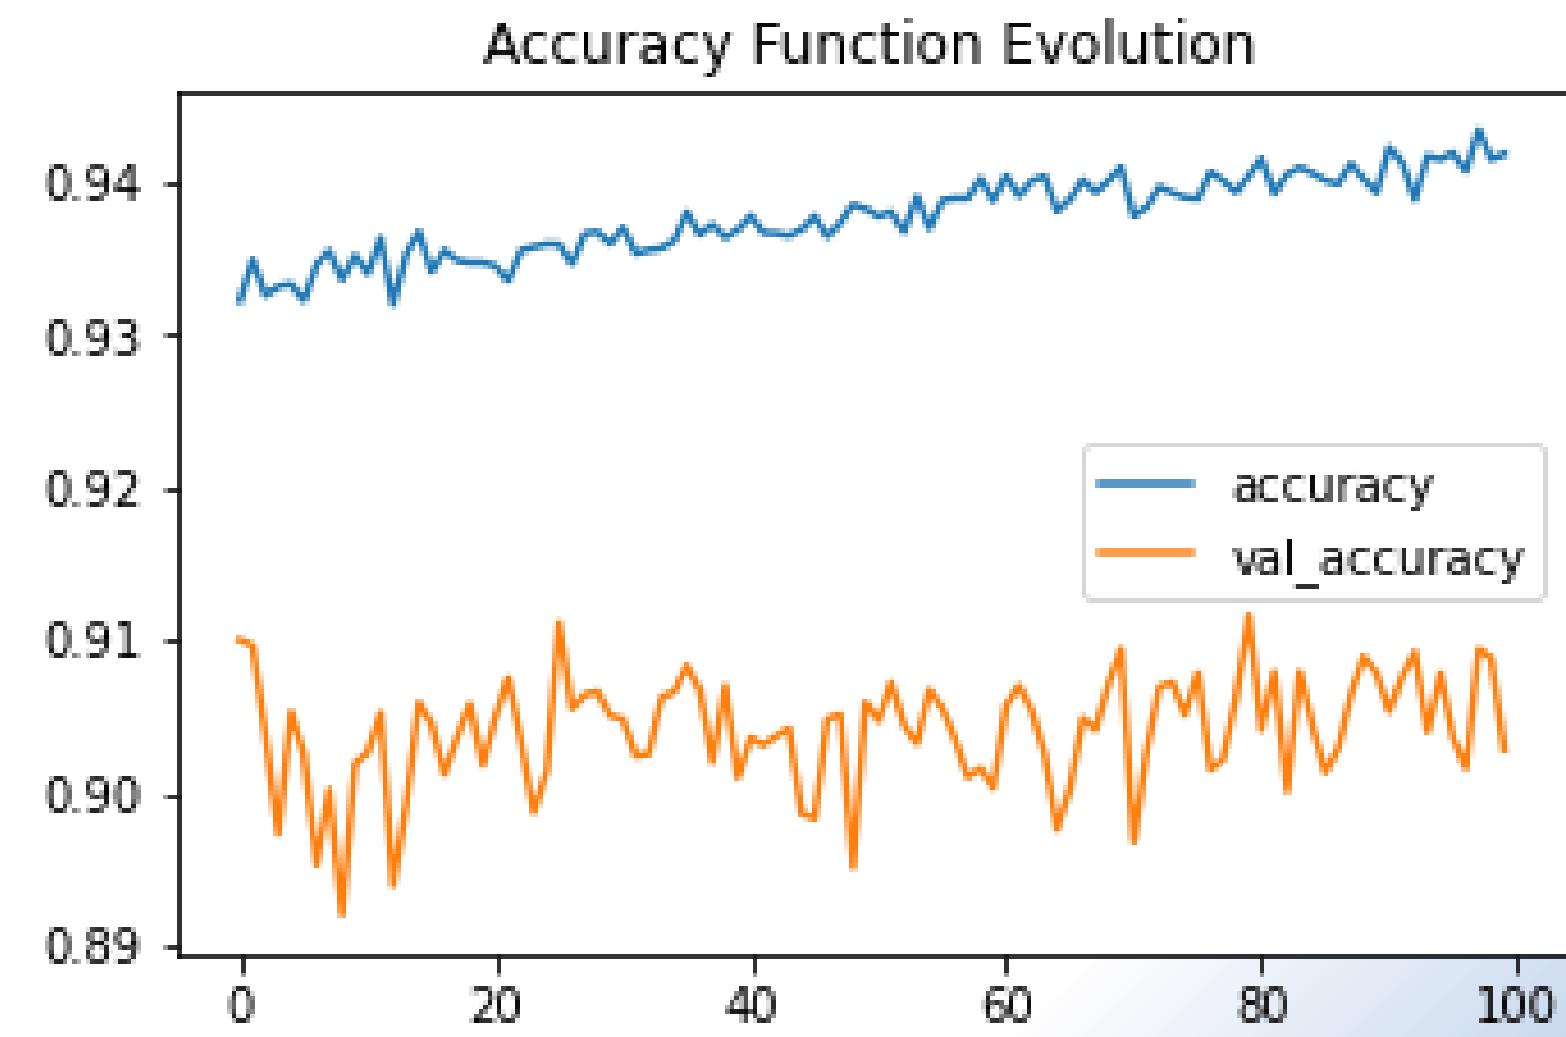
	parameter	Accuracy	loss
Optimizers	Adam	0.9418	0.0328
Loss function	binary cross entropy	0.9418	0.0328
Epochs	300	0.9418	0.0328
Batch size	64	0.9418	0.0328

< Comparision

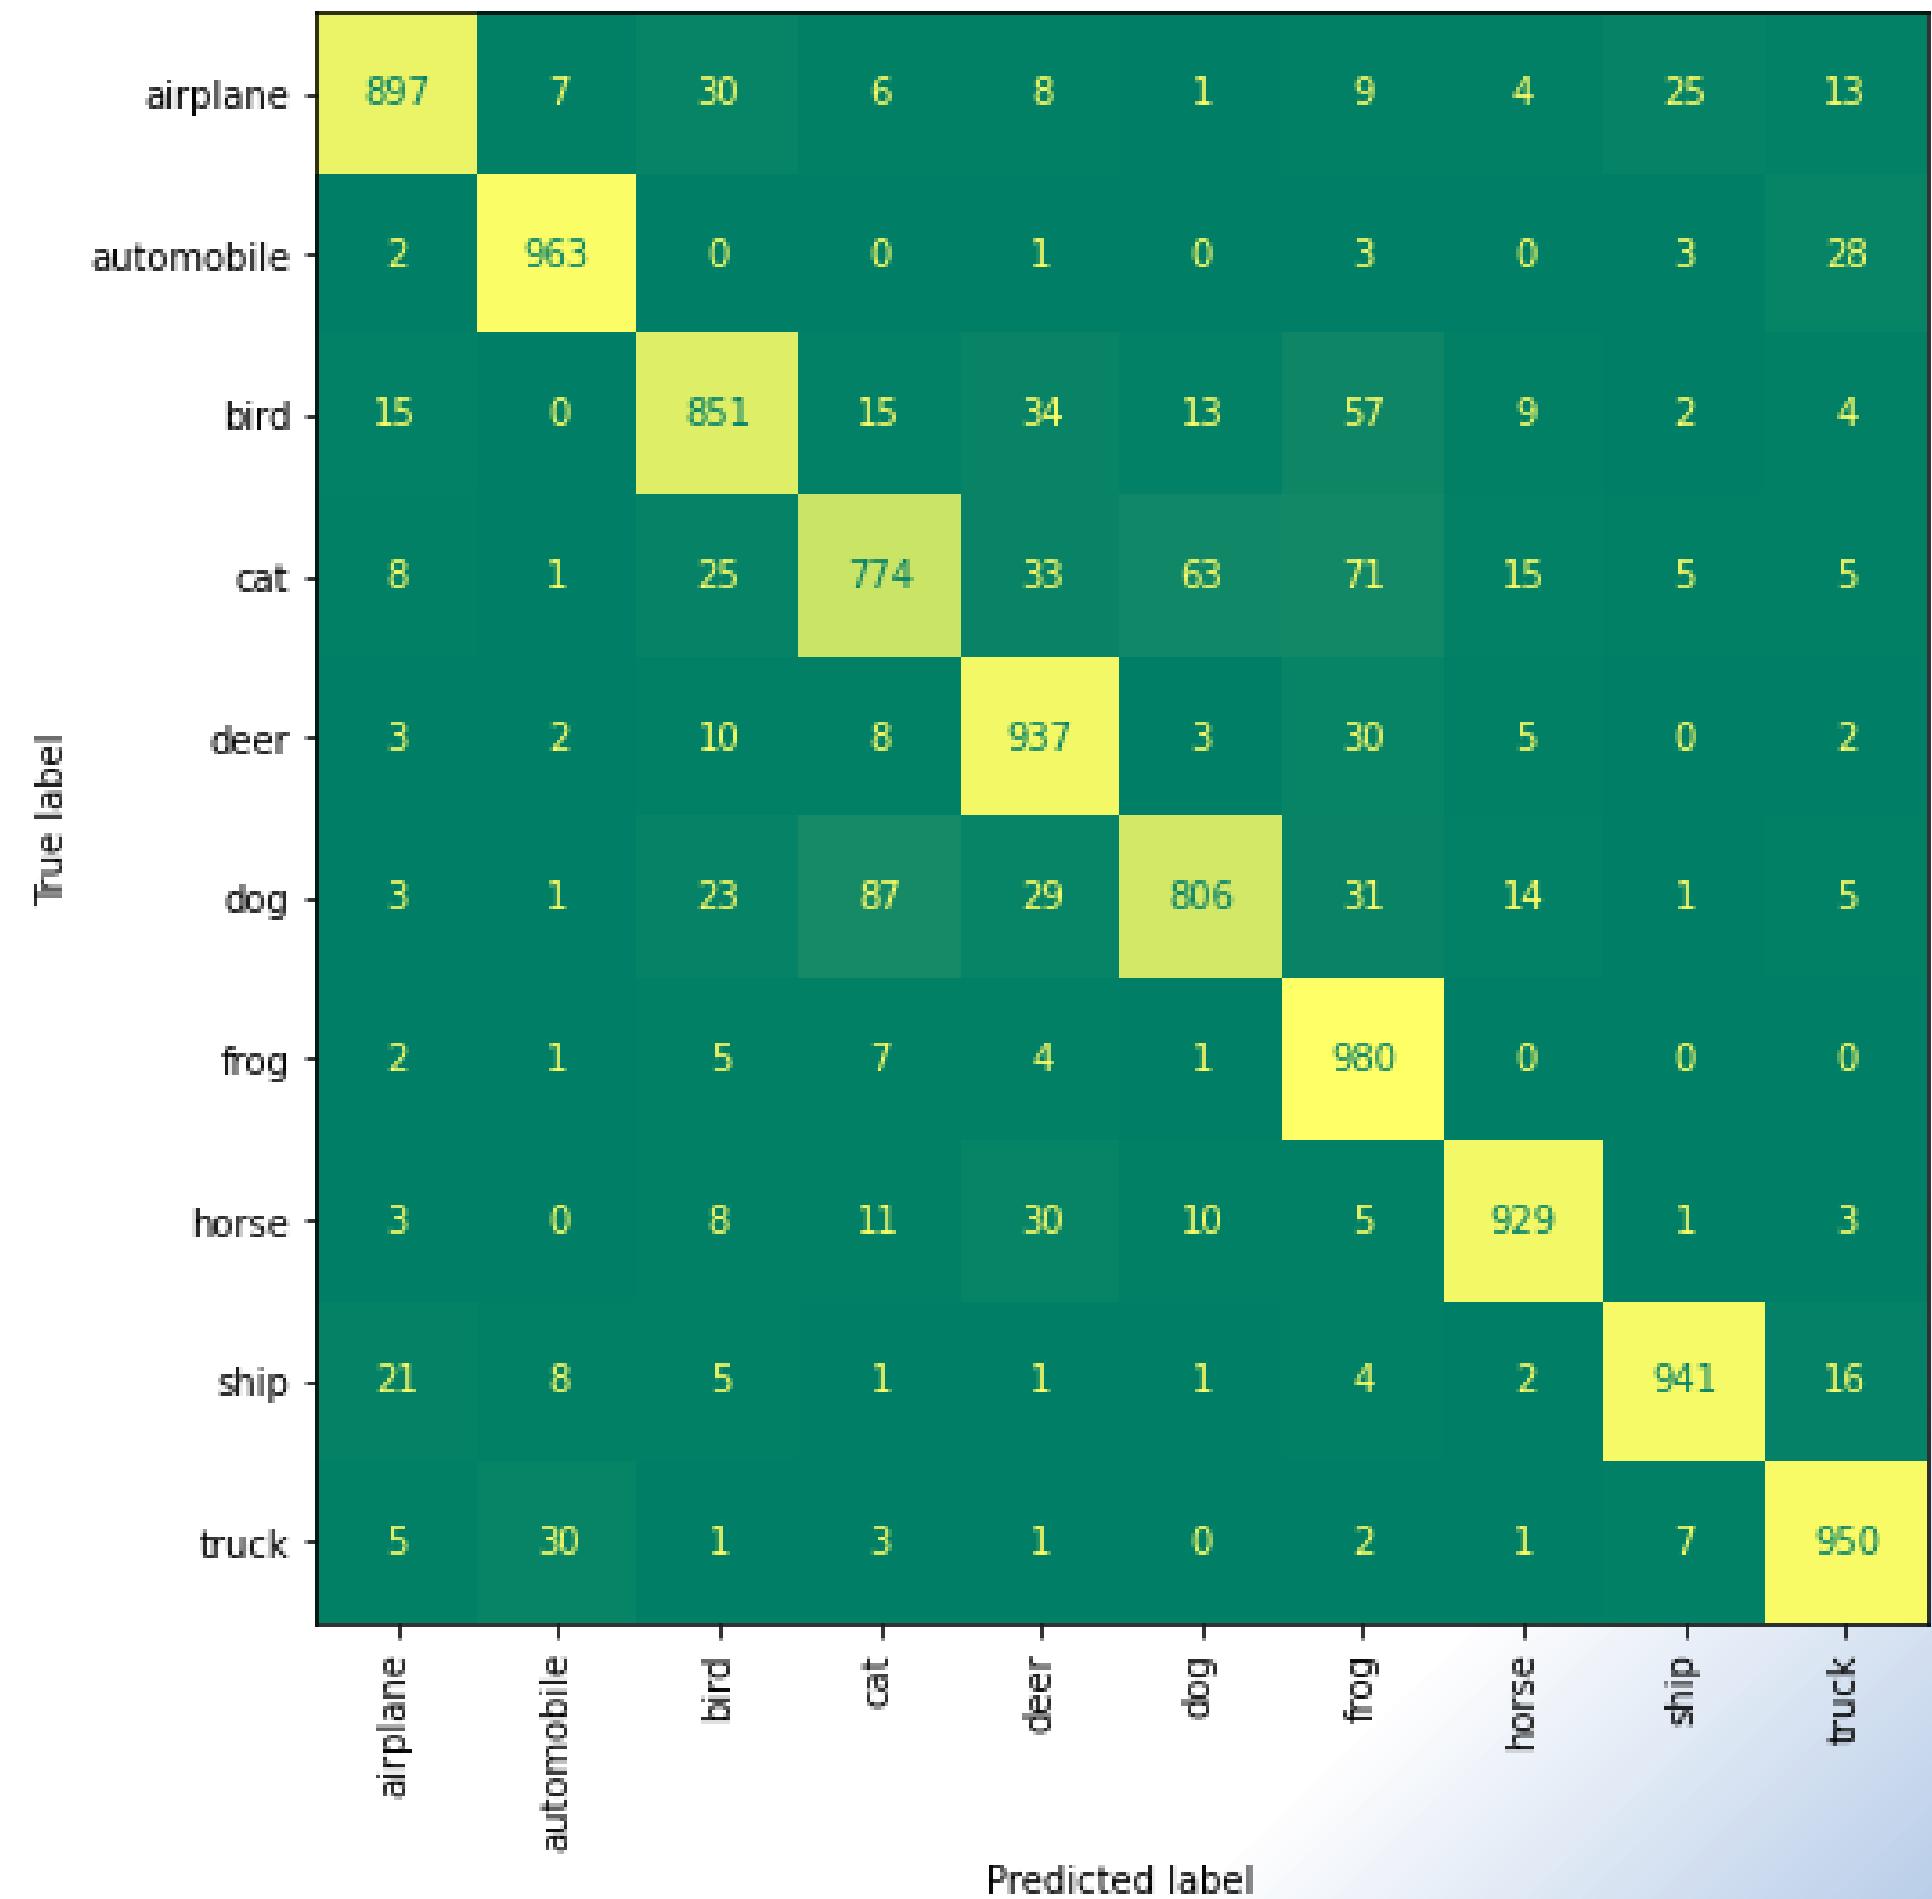
Loss curve



Accuracy curve



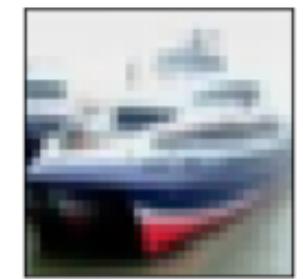
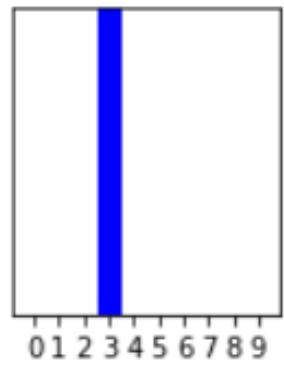
Heatmap



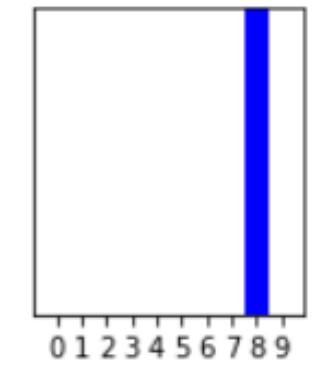
Comparision of Actual And predicted



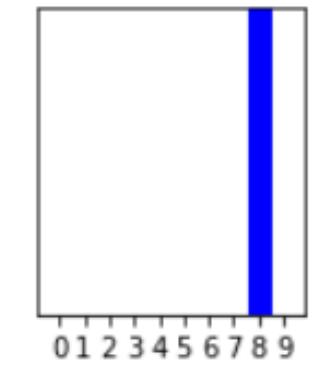
cat 100% (cat)



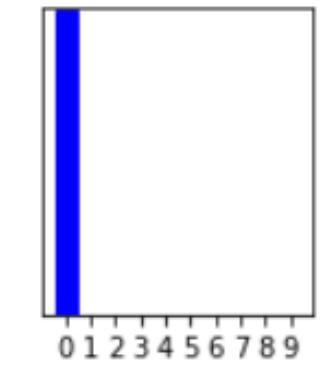
ship 100% (ship)



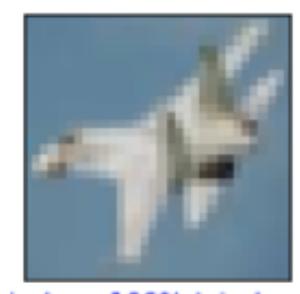
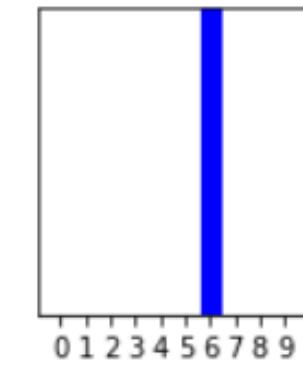
ship 100% (ship)



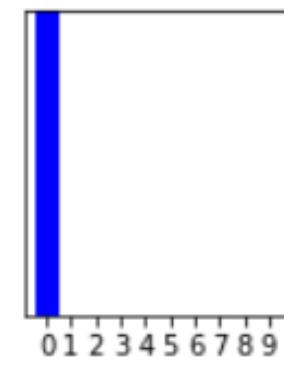
airplane 100% (airplane)



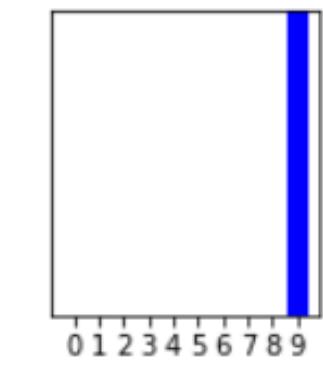
frog 100% (frog)



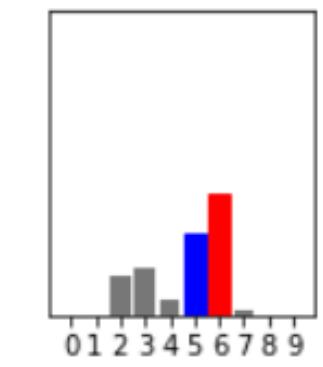
airplane 100% (airplane)



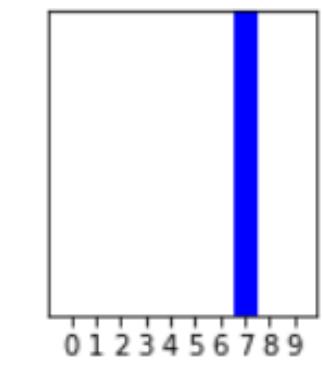
truck 100% (truck)



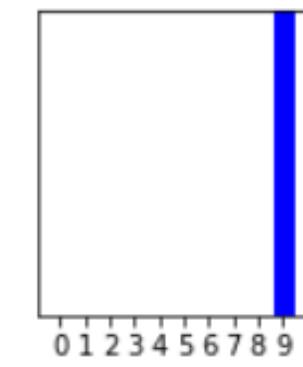
frog 40% (dog)



horse 100% (horse)



truck 100% (truck)



Conclusion:

- Optimizer- Adam
- Loss- Binary cross entropy
- Epochs- 300
- Batch size- 64

All the Images are equal in proportion in our dataset.

- Highest Accuracy-0.9418
- Least loss- 0.0328

Majority Correctly classified are Frog.

Minority correctly classified are cat.

Thank You



QR code of our code

APPENDIX

```
▶ import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
  
import tensorflow as tf  
from tensorflow.keras.datasets import cifar10  
from tensorflow.keras.utils import to_categorical  
  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization  
from tensorflow.keras.callbacks import EarlyStopping  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
  
from sklearn.metrics import ConfusionMatrixDisplay  
from sklearn.metrics import classification_report, confusion_matrix
```

```
▶ [(X_train, y_train)], (X_test, y_test) = cifar10.load_data()
```

```
👤 Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz  
170498071/170498071 [=====] - 13s 0us/step
```

```
[ ] print(f"X_train shape: {X_train.shape}")  
print(f"y_train shape: {y_train.shape}")  
print(f"X_test shape: {X_test.shape}")  
print(f"y_test shape: {y_test.shape}")
```

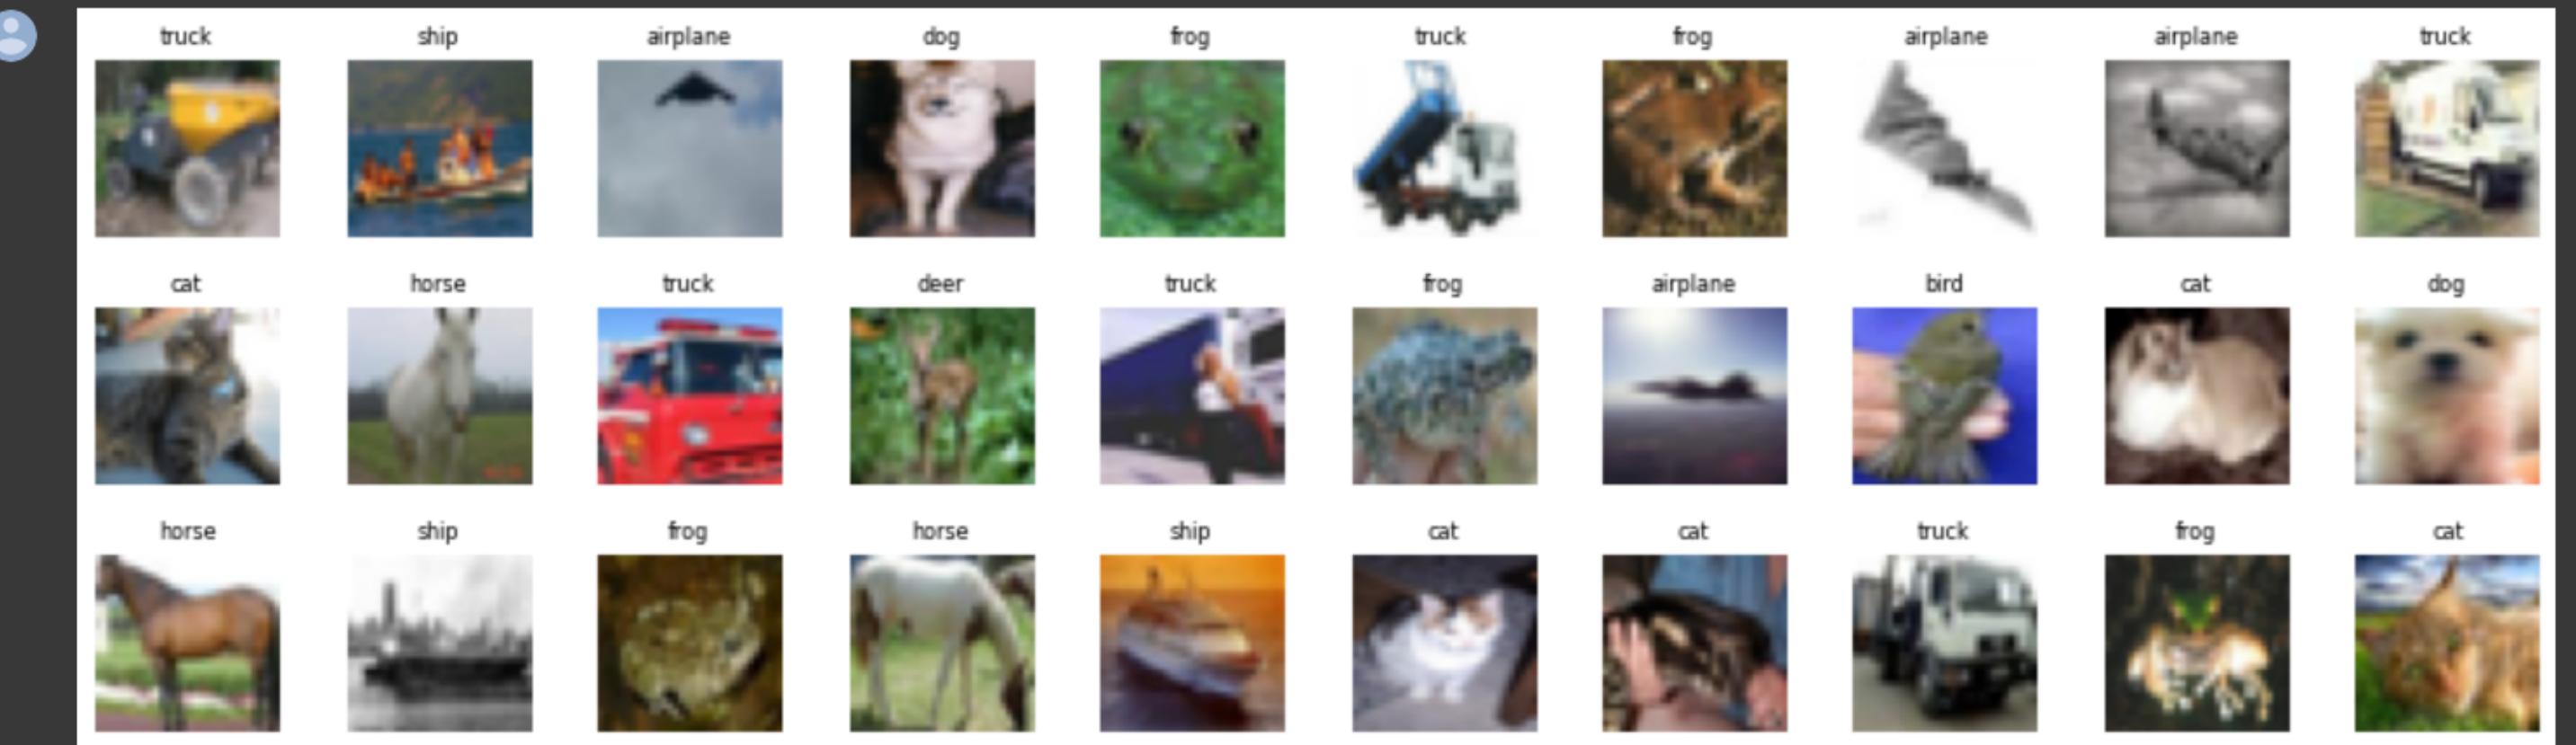
```
X_train shape: (50000, 32, 32, 3)  
y_train shape: (50000, 1)  
X_test shape: (10000, 32, 32, 3)  
y_test shape: (10000, 1)
```

```
▶ # Define the labels of the dataset  
labels = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
          'dog', 'frog', 'horse', 'ship', 'truck']  
  
# Let's view more images in a grid format  
# Define the dimensions of the plot grid  
W_grid = 10  
L_grid = 10  
  
# fig, axes = plt.subplots(L_grid, W_grid)  
# subplot return the figure object and axes object  
# we can use the axes object to plot specific figures at various locations  
  
fig, axes = plt.subplots(L_grid, W_grid, figsize = (17,17))  
  
axes = axes.ravel() # flatten the 15 x 15 matrix into 225 array
```

```
# Select a random number from 0 to n_train
for i in np.arange(0, W_grid * L_grid): # create evenly spaces variables

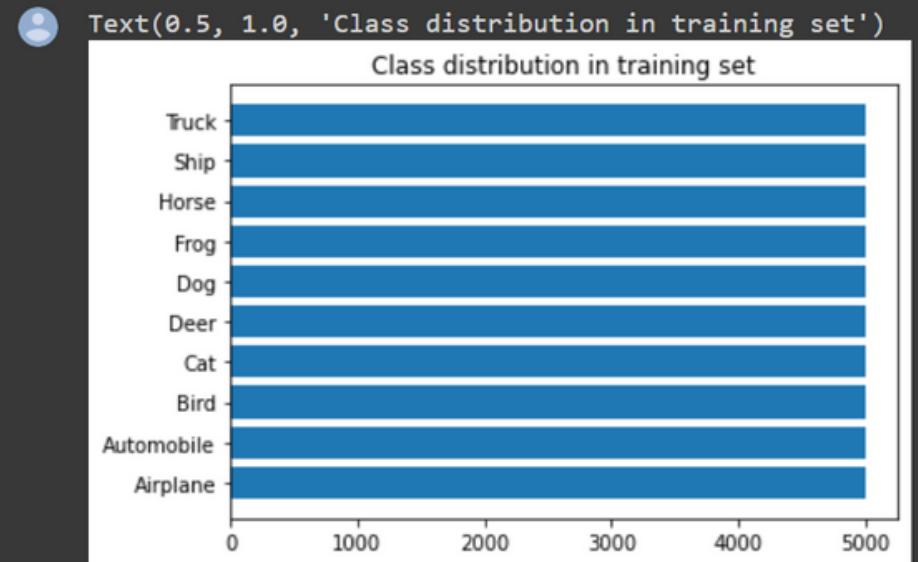
    # Select a random number
    index = np.random.randint(0, n_train)
    # read and display an image with the selected index
    axes[i].imshow(X_train[index,1:])
    label_index = int(y_train[index])
    axes[i].set_title(labels[label_index], fontsize = 8)
    axes[i].axis('off')

plt.subplots_adjust(hspace=0.4)
```

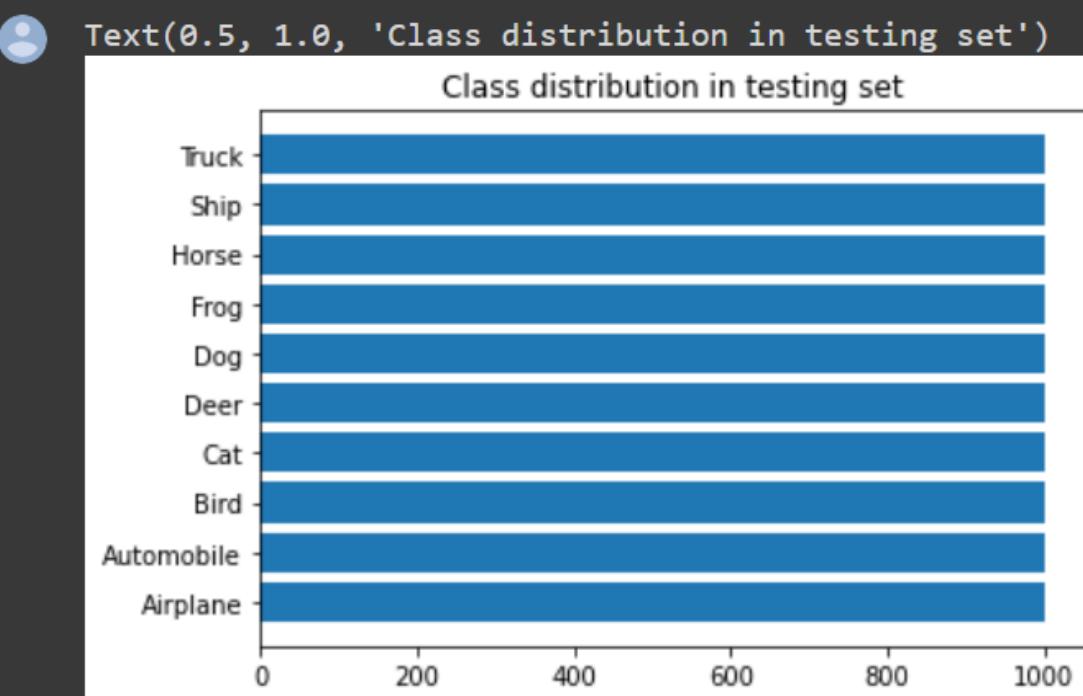


```
▶ classes_name = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse', 'Ship', 'Truck']

classes, counts = np.unique(y_train, return_counts=True)
plt.barh(classes_name, counts)
plt.title('Class distribution in training set')
```



```
▶ classes, counts = np.unique(y_test, return_counts=True)
plt.barh(classes_name, counts)
plt.title('Class distribution in testing set')
```



```
[ ] # Scale the data
X_train = X_train / 255.0
X_test = X_test / 255.0
```

```
y_cat_train  
array([[0., 0., 0., ..., 0., 0., 0.],  
       [0., 0., 0., ..., 0., 0., 1.],  
       [0., 0., 0., ..., 0., 0., 1.],  
       ...,  
       [0., 0., 0., ..., 0., 0., 1.],  
       [0., 1., 0., ..., 0., 0., 0.],  
       [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)  
  
[ ] model = Sequential()  
  
# Convolutional Layer  
model.add(Conv2D(filters=32, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(Conv2D(filters=32, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
# Pooling layer  
model.add(MaxPool2D(pool_size=(2, 2)))  
# Dropout layers  
model.add(Dropout(0.25))  
  
model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(MaxPool2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(Conv2D(filters=64, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(MaxPool2D(pool_size=(2, 2)))  
model.add(Dropout(0.25))  
  
model.add(Conv2D(filters=128, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(Conv2D(filters=128, kernel_size=(3, 3), input_shape=(32, 32, 3), activation='relu', padding='same'))  
model.add(BatchNormalization())  
model.add(MaxPool2D(pool_size=(2, 2)))  
model.add(Dropout(0.25))  
  
model.add(Flatten())  
# model.add(Dropout(0.2))  
model.add(Dense(128, activation='relu'))  
model.add(Dropout(0.25))  
model.add(Dense(10, activation='softmax'))  
  
METRICS = [  
    'accuracy',  
    tf.keras.metrics.Precision(name='precision'),  
    tf.keras.metrics.Recall(name='recall')  
]  
model.compile(loss='binary_crossentropy', optimizer='Adam', metrics=METRICS)
```

```
[ ] early_stop = EarlyStopping(monitor='val_loss', patience=2)

[ ] batch_size = 64
data_generator = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
train_generator = data_generator.flow(X_train, y_cat_train, batch_size)
steps_per_epoch = X_train.shape[0] // batch_size

▶ r = model.fit(train_generator,
                 epochs=100,
                 steps_per_epoch=steps_per_epoch,
                 validation_data=(X_test, y_cat_test),
                 #
                 callbacks=[early_stop],
                 batch_size=batch_size,
               )

❶ Epoch 1/100
781/781 [=====] - 26s 34ms/step - loss: 0.0373 - accuracy: 0.9322 - precision: 0.9431 - recall: 0.9237 - val_loss: 0.0569 - val_accuracy: 0.91
Epoch 2/100
781/781 [=====] - 26s 33ms/step - loss: 0.0360 - accuracy: 0.9349 - precision: 0.9452 - recall: 0.9266 - val_loss: 0.0558 - val_accuracy: 0.90
Epoch 3/100
781/781 [=====] - 26s 33ms/step - loss: 0.0371 - accuracy: 0.9325 - precision: 0.9436 - recall: 0.9247 - val_loss: 0.0598 - val_accuracy: 0.90
Epoch 4/100
781/781 [=====] - 26s 34ms/step - loss: 0.0369 - accuracy: 0.9331 - precision: 0.9438 - recall: 0.9245 - val_loss: 0.0639 - val_accuracy: 0.89
Epoch 5/100
781/781 [=====] - 26s 33ms/step - loss: 0.0367 - accuracy: 0.9333 - precision: 0.9448 - recall: 0.9249 - val_loss: 0.0585 - val_accuracy: 0.90
Epoch 6/100
781/781 [=====] - 27s 34ms/step - loss: 0.0372 - accuracy: 0.9322 - precision: 0.9434 - recall: 0.9238 - val_loss: 0.0616 - val_accuracy: 0.90
Epoch 7/100
781/781 [=====] - 26s 33ms/step - loss: 0.0363 - accuracy: 0.9346 - precision: 0.9459 - recall: 0.9267 - val_loss: 0.0657 - val_accuracy: 0.89
Epoch 8/100
```

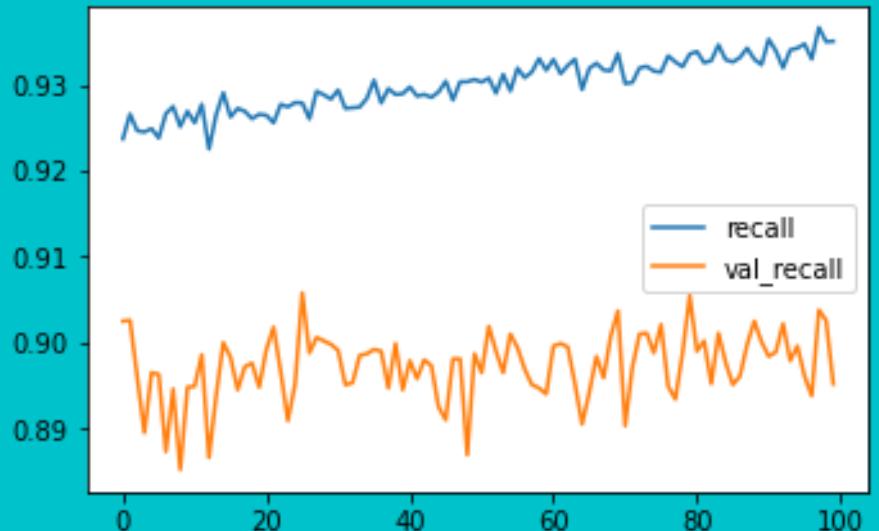
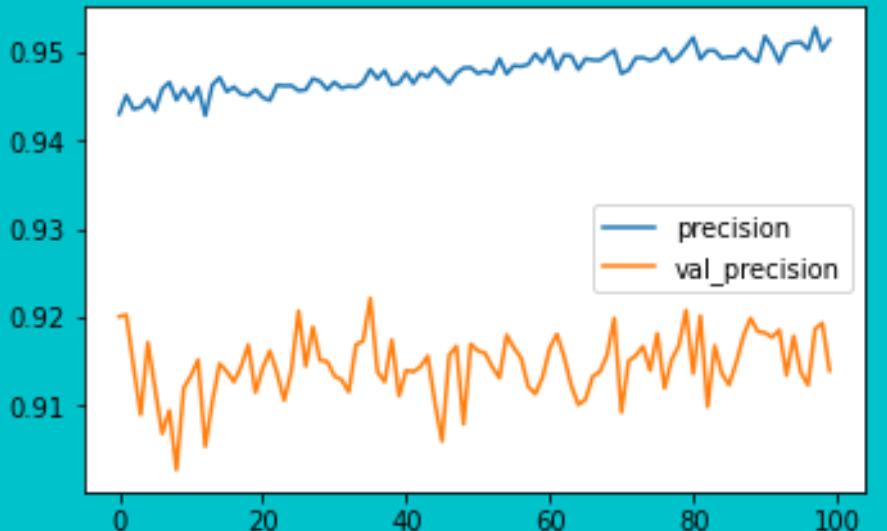
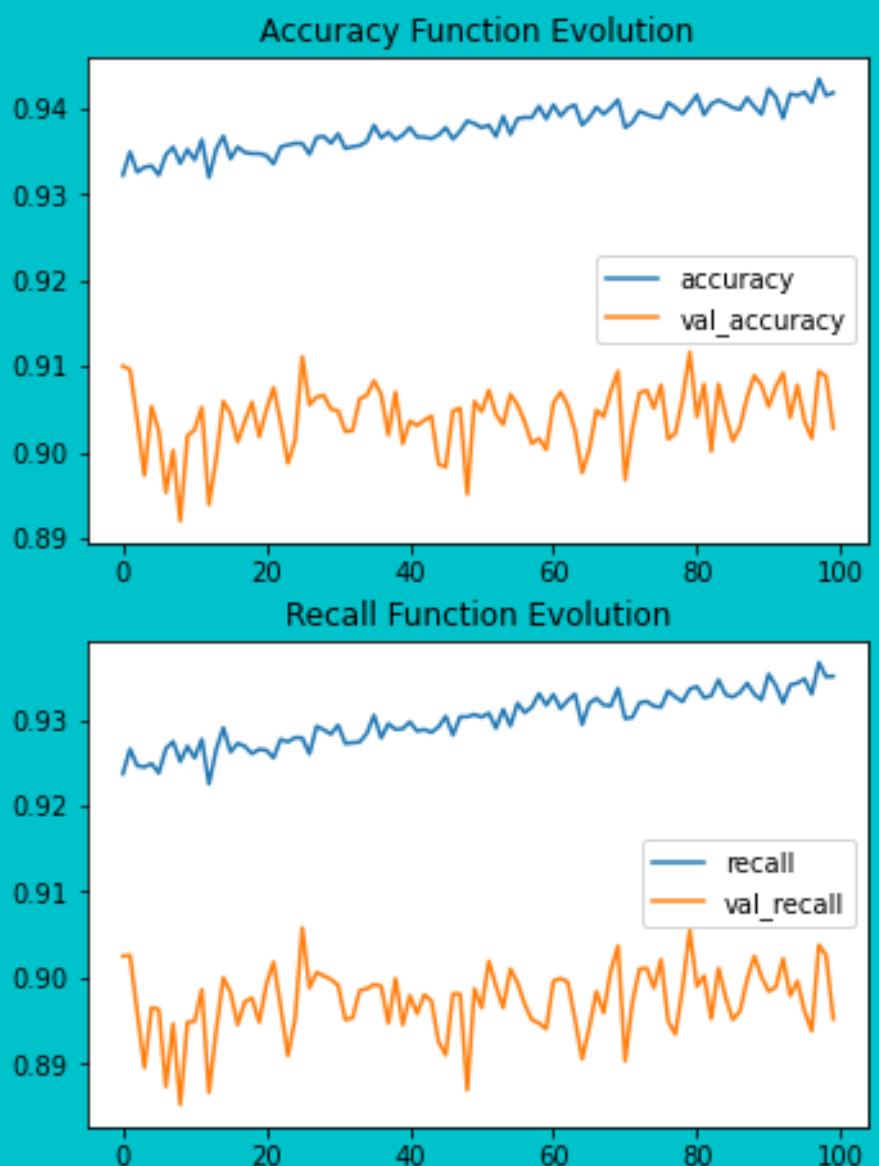
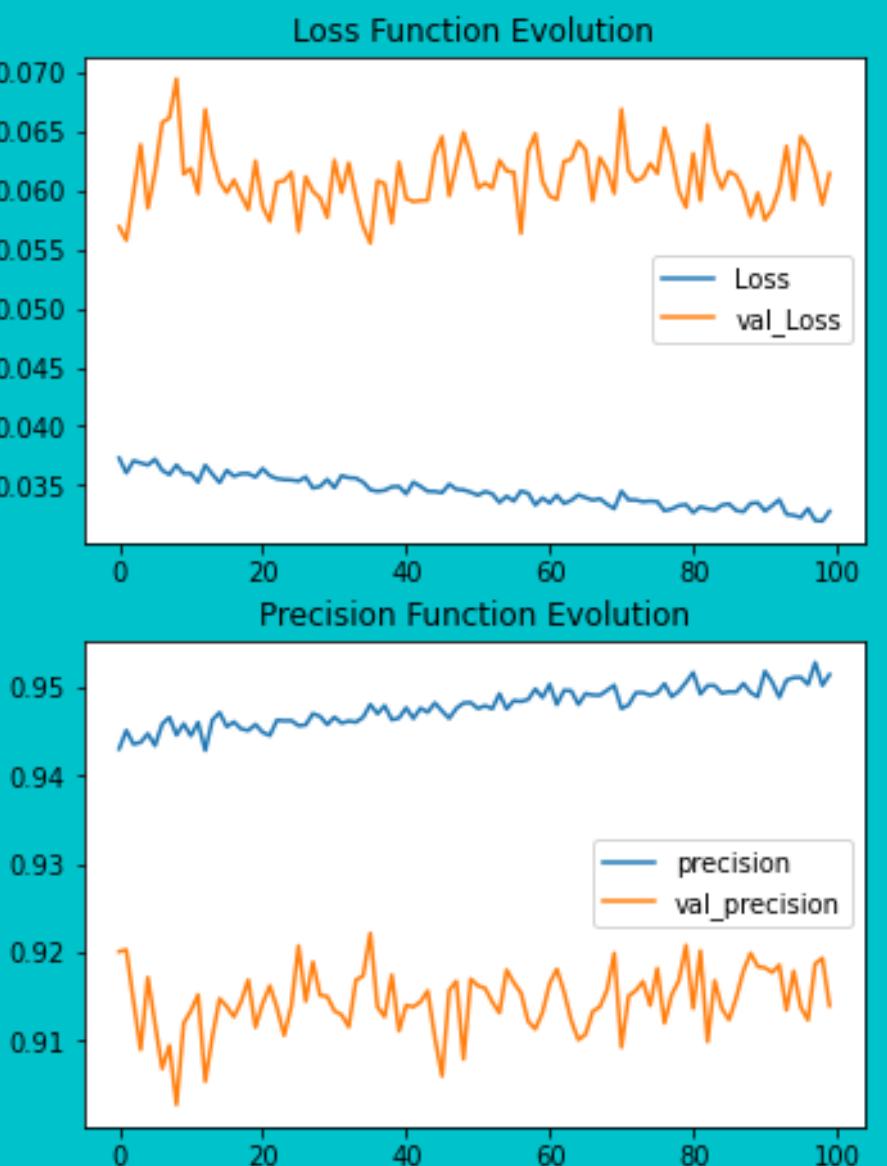
```
from matplotlib import pyplot as plt
plt.figure(figsize=(12, 16))

plt.subplot(4, 2, 1)
plt.plot(r.history['loss'], label='Loss')
plt.plot(r.history['val_loss'], label='val_Loss')
plt.title('Loss Function Evolution')
plt.legend()

plt.subplot(4, 2, 2)
plt.plot(r.history['accuracy'], label='accuracy')
plt.plot(r.history['val_accuracy'], label='val_accuracy')
plt.title('Accuracy Function Evolution')
plt.legend()

plt.subplot(4, 2, 3)
plt.plot(r.history['precision'], label='precision')
plt.plot(r.history['val_precision'], label='val_precision')
plt.title('Precision Function Evolution')
plt.legend()

plt.subplot(4, 2, 4)
plt.plot(r.history['recall'], label='recall')
plt.plot(r.history['val_recall'], label='val_recall')
plt.title('Recall Function Evolution')
plt.legend()
```



```
[ ] evaluation = model.evaluate(X_test, y_cat_test)
print(f'Test Accuracy : {evaluation[1] * 100:.2f}%')

313/313 [=====] - 1s 4ms/step - loss: 0.0614 - accuracy: 0.9028 - precision: 0.9139 - recall: 0.8952
Test Accuracy : 90.28%

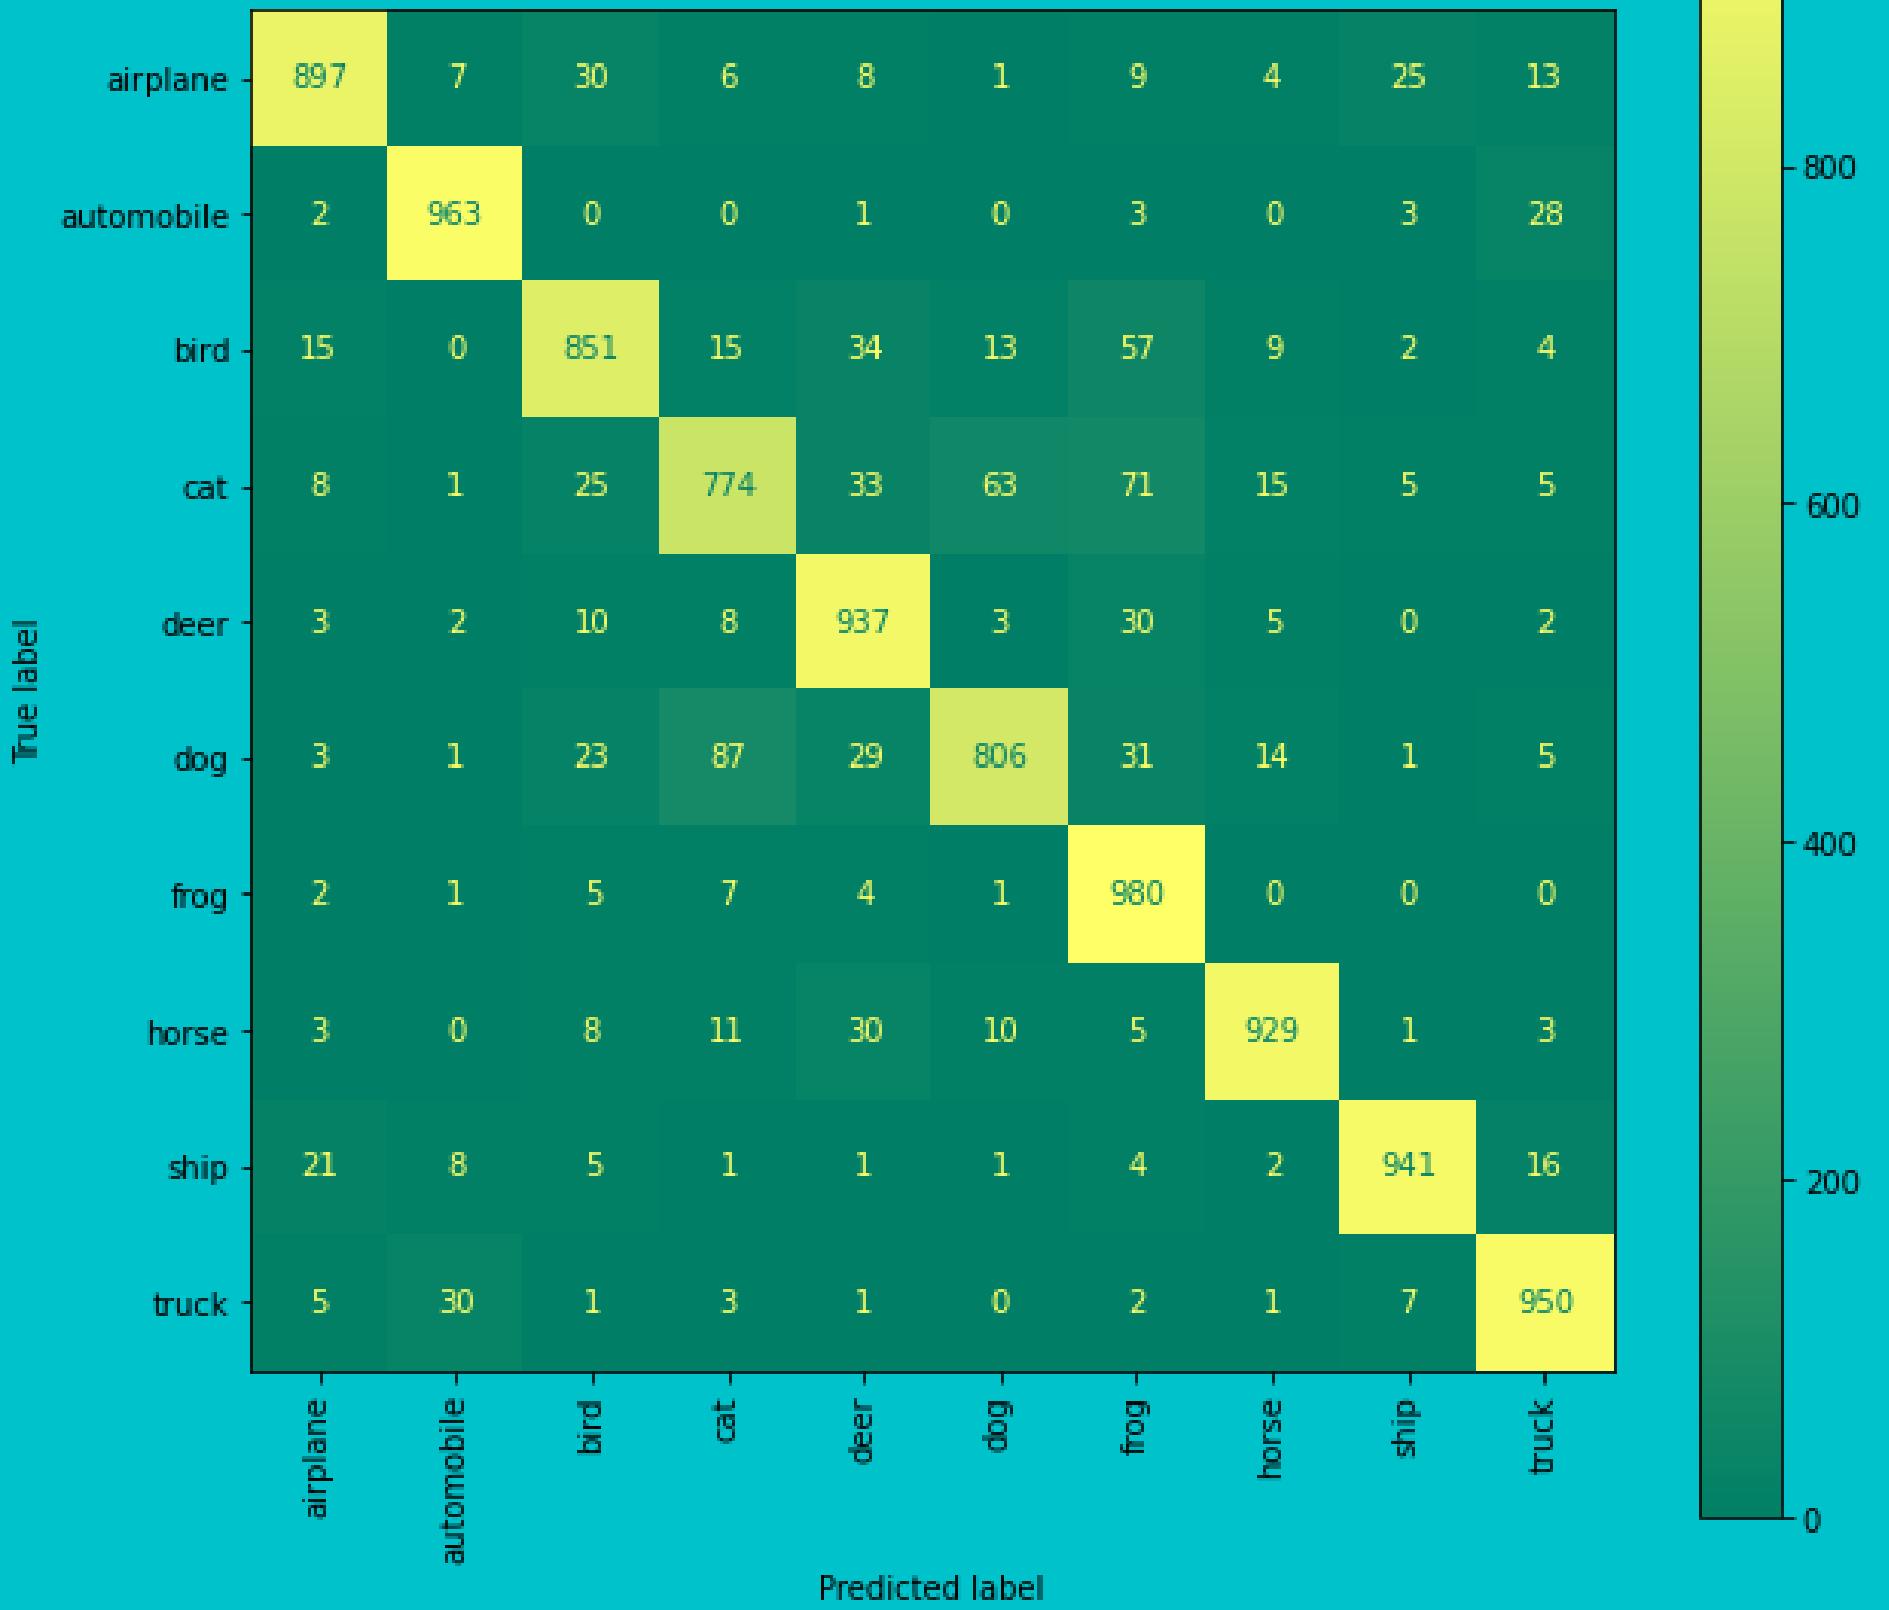

[ ] y_pred = model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)
cm = confusion_matrix(y_test, y_pred)

313/313 [=====] - 1s 3ms/step


[ ] disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                                 display_labels=labels)

# NOTE: Fill all variables here with default values of the plot_confusion_matrix
fig, ax = plt.subplots(figsize=(10, 10))
disp = disp.plot(xticks_rotation='vertical', ax=ax, cmap='summer')

plt.show()
```



```
[ ] print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	1000
1	0.95	0.96	0.96	1000
2	0.89	0.85	0.87	1000
3	0.85	0.77	0.81	1000
4	0.87	0.94	0.90	1000
5	0.90	0.81	0.85	1000
6	0.82	0.98	0.89	1000
7	0.95	0.93	0.94	1000
8	0.96	0.94	0.95	1000
9	0.93	0.95	0.94	1000
accuracy			0.90	10000
macro avg	0.90	0.90	0.90	10000
weighted avg	0.90	0.90	0.90	10000

```
▶ def plot_image(i, predictions_array, true_label, img):
    predictions_array, true_label, img = predictions_array, true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img, cmap=plt.cm.binary)

    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'

    plt.xlabel(f'{labels[int(predicted_label)]} {100*np.max(predictions_array):2.0f}% ({labels[int(true_label)]})',
               color=color)

def plot_value_array(i, predictions_array, true_label):
    predictions_array, true_label = predictions_array, int(true_label[i])
    plt.grid(False)
    plt.xticks(range(10))
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(predictions_array)

    thisplot[predicted_label].set_color('red')

    thisplot[predicted_label].set_color('red')
    thisplot[true_label].set_color('blue')
```

```
▶ # Plot the first X test images, their predicted labels, and the true labels.
# Color correct predictions in blue and incorrect predictions in red.

num_rows = 8
num_cols = 5
num_images = num_rows * num_cols
plt.figure(figsize=(2 * 2 * num_cols, 2 * num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2 * num_cols, 2 * i + 1)
    plot_image(i, predictions[i], y_test, X_test)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(i, predictions[i], y_test)
plt.tight_layout()
plt.show()
```



